

Contents

1	Meeting with Murali's students	1
1.1	Planning to finish the article [0/4]	1
1.1.1	new stuff over previous paper	1
1.1.2	TODO experiments [0/2]	2
1.1.3	TODO missing proofs	2
1.1.4	TODO replace latex template	2
1.1.5	TODO writing [0/6]	2
1.1.6	questions	3
1.2	TODO what to introduce everybody before	3
1.2.1	TODO databases	3
1.2.2	union in datalog	4
1.2.3	DONE terminal basics	12
1.2.4	TODO programming languages	13
1.3	setup access to machines	13
1.3.1	machines	13
1.3.2	ssh	13
1.3.3	running postgres	14
1.4	code and what to compare	15
1.5	experiments	15
1.5.1	the problem setting	15
1.5.2	what parameters to vary?	15
1.5.3	competitors	15
1.5.4	running step 2	16
1.5.5	running step 3	16
1.5.6	TODO running things in PDQ	19
1.5.7	experiment dry run	19
1.6	TPC-H datalog queries	19
1.6.1	Q3	19
1.6.2	Q18	20
2	Meeting Murali	20

1 Meeting with Murali's students

Added: *[2021-12-29 Wed 12:38]*

1.1 Planning to finish the article [0/4]

- VLDB Journal

1.1.1 new stuff over previous paper

- new algorithm with variants
- new experiments
 - comparison with PDQ

– optimization time + execution time breakdown

1.1.2 TODO experiments [0/2]

☐ TPC-H

- ☐ translate queries into Datalog (8 / 19?) (**4 hours**)
- ☐ measure optimization time and runtimes (**4 hours**)
- ☐ measure PDQ optimization times (**16 hours**)
- ☐ analyze and plot (**24 hours**)

☐ real world data

- ☐ select datasets
- ☐ translate all queries into Datalog
- ☐ measure optimization time and runtimes
- ☐ measure PDQ optimization times
- ☐ analyze and plot

1.1.3 TODO missing proofs

1.1.4 TODO replace latex template

1.1.5 TODO writing [0/6]

1. TODO introduction [1/4]

- ☐ add running example
 - ☐ hopefully greedy does not work (if we can keep it simple enough. Otherwise, this example comes later)
- ☒ prior (new results of IPAW paper)
- ☐ contributions list
- ☐ outline of paper

2. TODO background [1/4]

- ☐ Datalog
- ☐ provenance and capture
- ☒ dependencies
- ☐ semantic query optimization with chase & backchase

3. TODO related work [0/2]

- ☐ provenance capture & querying
 - ☐ check for new papers
- ☐ semantic query optimization
 - ☐ more citations including (provenance-directed) chase & backchase, PDQ-related papers, ...

4. **TODO** restricted provenance capture (capture for one input table) [0/2]

- ☐ explain basic idea
- ☐ write rules

5. **TODO** semantic optimization algorithm [0/3]

- ☐ write algorithm
- ☐ explain algorithm
- ☐ proof correctness

6. **TODO** write experiments [0/4]

- ☐ explain setup and competitors
- ☐ optimization time comparison with PDQ
- ☐ runtime comparison with PDQ and unoptimized version
- ☐ datasets & workloads
 - ☐ TPC-H
 - ☐ real world data

1.1.6 questions

- where would the algorithms show benefit over the old ones
-

1.2 TODO what to introduce everybody before

1.2.1 TODO databases

1. **TODO** SQL
2. **TODO** administrating a database (Postgres)
3. **TODO** datalog

- Table: **Person**

Name	Age	Salary	Dept
Peter	34	3444	CS
Bob	60	10500	CS
Alice	34	40000	CS
Fred	3444	10	News

- Table: **Department**

Title	Budget	City
CS	134300	Chicago
News	123	Ann Arbor

- active domain $adom(D)$ is all values that exist the database

```

Q(name,salary) :- Person(name,_,salary,"CS").

-- name: Peter, age: 60, salary: 40000
Q(Peter,40000) :- Person(Peter,60,40000,"CS").
-- name: Peter, age: 34, salary: 40000
Q(Peter,40000) :- Person(Peter,34,40000,"CS").
-- name: Peter, age: 34, salary: 3444
Q(Peter,3444) :- Person(Peter,34,3444,"CS"). -- works out
-- name: Peter, age: 60, salary: 3444
Q(Peter,3444) :- Person(Peter,_,3444,"CS"). -- works out
-- name: 60, age: 60, salary: 60

```

```

CREATE TABLE person (
  name VARCHAR(15) PRIMARY KEY,
  age INT,
  salary INT,
  dept VARCHAR(10)
);

```

```

Q(Name) :- Person(Name,X,Y,Z).

```

- get all employee's from the CS department

```

Q(name,salary,dept) :- Person(name,_,salary,dept), dept="CS".
Q(name,salary,dept) :- Person(name,_,salary,"CS").

```

- get all employees that earn more then 100000

```

Q(name,salary,dept) :- Person(name,_,salary,dept), salary > 100000.

Q(Name,Salary,Dept) :- person(Name,A,Salary,Dept), department(Dept,B,C).
Q(X,Y,Z) :- person(X,_,Y,Z), department(Z,B,C).

```

1.2.2 union in datalog

Person

Name	Dept
Peter	CS
Alice	CS
Bob	HR
Gert	HR

```

Q(Name) :- Person(Name, cs).
Q(Name) :- Person(Name, hr).

```

1. TODO equivalence, constraints, semantic query optimization

```

Q(X) :- R(X,Y). -- R is edb
Q2(X) :- Q(X). -- Q + Q2 is idb

```

- <http://www.cs.iit.edu/~glavic/cs520/2022-spring/exams/>

(a) Set notation $S = \{e_1, \dots, e_n\}$

$$S_1 = \{a, b, c\} \text{ and } S_2 = \{a, b\}$$

$$S_1 \supseteq S_2$$

$$S_1 \supset S_2$$

$$S_1 = S_2 \Leftrightarrow S_1 \subseteq S_2 \wedge S_1 \supseteq S_2$$

(b) Predicate and First-order Logic

- AND \wedge , OR \vee , NOT \neg , implies \rightarrow

x	y	x AND y
0	0	0
0	1	0
1	0	0
1	1	1

x	y	x \rightarrow y
0	0	1
0	1	1
1	0	0
1	1	1

- variables over domain values \mathbb{N}
- predicates $<: \mathbb{N} \times \mathbb{N} \rightarrow \{F, T\}$
 - $x < y, x < y \wedge y < z$
- quantification:
 - $\forall x : \phi(x)$ - is true if for all x the $\phi(x)$ (universal)
 - * $\forall x : bird(x) \rightarrow canfly(x)$

if x is bird then x canfly? probably not true

$$\forall x : isCSstudent(x) \rightarrow canprogram(x)$$

(c) Equivalence and Containment

- Queries Q are function:
 - Input: database D (EDB only)
 - Output: database $Q(D)$ (EDB + IDB)
- **Query equivalence:**
 - Q and Q' are equivalent if for every database D we have $Q(D) = Q'(D)$

$$Q \equiv Q' \Leftrightarrow \forall D : Q(D) = Q'(D)$$

- **Query containment:**
 - Q is contained in Q' if for every database D we have $Q(D)$ subset or equal to $Q'(D)$

$$Q \sqsubseteq Q' \Leftrightarrow \forall D : Q(D) \subseteq Q'(D)$$

- Write query equivalence as containment

$$Q \equiv Q' \Leftrightarrow \forall D : Q(D) \subseteq Q'(D) \wedge Q'(D) \subseteq Q(D) \Leftrightarrow Q \sqsubseteq Q' \wedge Q' \sqsubseteq Q$$

```

Q1(X) :- R(X,Y), R(X,Z).
Q2(X) :- R(X,Y).
Q3(X) :- Q2(X), X < 2. -- Dominick Q3 contained in Q2
Q4(Y) :- R(X,Y). -- not contained in Q2
Q5(X) :- R(X,Y).
Q6(A) :- R(A,GFDFGDFG). -- equivalent to Q2 and Q5
Q7(X) :- R(X,Y), X < 2.
Q8(X) :- R(X,Y), S(Y,Z). -- contained Q2

```

- **Variable names are irrelevant:** only their positions in the body and head matter
- **One body is superset of another body:** it is more restrictive (it returns less results)
- **if two queries do not return the "same" variables (after renaming):** no containment relationship
- **Containment mapping**
 - Variable mapping $\text{Var}(Q) \rightarrow \text{Var}(Q')$ then we rename all variables from Q to variables from Q'
 - * $Q2 \rightarrow Q5$: $X \rightarrow X, Y \rightarrow X, X \rightarrow X, Y \rightarrow Y, X \rightarrow Y, Y \rightarrow Y, X \rightarrow Y, Y \rightarrow X$
 - Containment mapping is a variable mapping that fulfills these two conditions:
 - * 1) the head is head mapped to the head
 - * 2) every atom from the body of Q exists after renaming in the body of Q'
- **Example:**
 - $Q2 \rightarrow Q6$:
 - * head to head: $X \rightarrow A$
 - * body to body: $Y \rightarrow \text{GFDFGDFG}$ we get $R(X,Y)$ is mapped to $R(A,\text{GFDFGDFG})$
 - $Q2 \rightarrow Q5$:
 - * head to head: $X \rightarrow X$
 - * body to body: $Y \rightarrow Y$
 - $Q7 \rightarrow Q2$
 - * head to head: $X \rightarrow X$
 - * body to body: $Y \rightarrow Y$
 - $Q8 \rightarrow Q2$
 - * $X \rightarrow X$
 - * $Y \rightarrow Y$
 - * $=Z \rightarrow =$
 - * head to head: YES
 - * body to body: NO
 - $Q2 \rightarrow Q8 \rightarrow Q_8 \sqsubseteq Q_2$
 - * $X \rightarrow X$
 - * $Y \rightarrow Y$
 - * head to head: YES

* body to body: YES
 – Q1 \rightarrow Q2 \rightarrow $Q_2 \sqsubseteq Q_1$
 * X \rightarrow X
 * Y \rightarrow Y
 * Z \rightarrow Y
 * head to head: YES
 * body to body: YES
 – Q2 \rightarrow Q1 \rightarrow $Q_1 \sqsubseteq Q_2$
 * X \rightarrow X
 * Y \rightarrow Y
 * head to head: YES
 * body to body: YES

R

A	B
a	1
b	2
c	3

Q1(D)

X
a
b
c

Q2(D)

X
c
a
b

2. TODO constraints

- **primary key**: attributes of a table that are unique in a table
- **SSN** as PK for this table

SSN	Name	Salary
111	Peter	30003
222	Peter	12312
333	Bob	12312

(a) functional dependencies

- **functional dependencies**
 - A \rightarrow B holding over R
 - then for any two tuple $t, t' \in R$ if $t.A = t'.A$ then $t.B = t'.B$
 - SSN \rightarrow Name, Salary

Name	Zip	city
Peter	60616	Chicago
Bob	60616	Chicago
Alice	60657	Chicago
Fred	11111	New York

- evaluate query under the knowledge that `zip -> city` holds for the database

$Q(C1, C2) :- \text{address}(_, Z, C1), \text{address}(_, Z, C2), C1 \neq C2.$

- result is guaranteed to be empty when know that `zip -> city` holds

(b) fulfilling constraints

Name	zip	city
Peter	60616	Chicago
Bob	60616	New York
Alice	60657	Chicago
Fred	11111	New York

(c) inclusion dependencies

- foreign keys as a special case
- `Person(Name, LiveAt), Address(Id, City, Zip, Street)` with `Id` is PK for `address`
- `Person(Name, Id, City, Zip, Street)` is also an option

Person

Name	LivesAt
Peter	1
Alice	1
bob	2

Address

Id	city	zip	street
1	Chicago	60614	adsasdas
2	Chicago	60666	adsasdasdasd
3	Chicago	60615	adsasdalosjkdas

- foreign key constraint. For every value of attribute **A** of table **R** there has to exists tuple **s** in table **S** with PK equal to the value of **A**.
 - the set of values in attribute **LivesAt** has to be a subset of the values in attribute **Id**
 - **inclusion dependency**

$$\forall \text{name, livesat} : \text{Person}(\text{name}, \text{livesat}) \rightarrow \exists \text{city, zip, street} : \text{Address}(\text{livesat}, \text{city}, \text{zip}, \text{street})$$

3. **TODO** semantic query optimization

(a) Semantic query optimization problem

- **Inputs:**
 - database **D** and set constraint Σ
 - query **Q**

- **Output:**

- query Q' that is equivalent to Q under the Σ
- "optimal in some way"

(b) Example

- `Person(Name,LiveAt), Address(Id,City,Zip,Street)` with `Id` is PK for address

`Q1(N) :- Person(N,L), Address(L,C,Z,S).`

`Q2(N) :- Person(N,L).`

Person

Name	LivesAt
Peter	3
Alice	1
bob	2

Address

Id	city	zip	street
1	Chicago	60614	adsasdas
2	Chicago	60666	adsasdasdasd
3	Chicago	60615	adsasdalosjkdas

(c) Query minimization

- find smallest query Q' such that $body(Q') \subseteq body(Q)$ and that $Q \equiv Q'$
 - size of Q is measured as number of atoms in the body of Q
 - we have function `equivalent(Q,Q') -> Bool` and have function `unsafe(Q) -> Bool`

`Q(N) :- Person(N,L), Address(L,C,Z,S).`

`Q1(N) :- Address(L,C,Z,S). -- unsafe`

`Q2(N) :- Person(N,L).`

`Q3(N) :- . -- unsafe`

`Q(X,Y) :- R(X,Y), R(X,Z), R(X,A).`

- $S = \{a,b,c\}, \dots, \{\}, \{a\}, \{b\}, \{c\}, \{a,b\}, \{a,c\}, \{b,c\} \dots$

`Q1(X,Y) :- R(X,Y), R(X,Z). -- safe, equivalent`

`Q2(X,Y) :- R(X,Y), R(X,A). -- safe, equivalent`

`Q3(X,Y) :- R(X,Z), R(X,A). -- unsafe`

`Q4(X,Y) :- R(X,Y). -- safe, equivalent`

`Q5(X,Y) :- R(X,A). -- unsafe`

`Q6(X,Y) :- R(X,Z). -- unsafe`

`Q7(X,Y) :- . -- unsafe`

- Equivalence of Q and $Q1$, $Q \equiv Q' \Leftrightarrow Q \sqsubseteq Q' \wedge Q' \sqsubseteq Q$

`Q(X,Y) :- R(X,Y), R(X,Z), R(X,A).`

`Q1(X,Y) :- R(X,Y), R(X,Z).`

- $Q \rightarrow Q1$:

- CM: $X \rightarrow X, Y \rightarrow Y, Z \rightarrow Z, A \rightarrow Y$

- $Q1 \rightarrow Q$

– CM: $X \rightarrow X, Y \rightarrow Y, Z \rightarrow Y$

- revisiting person example

$Q(N) :- \text{Person}(N,L), \text{Address}(L,C,Z,S).$

$Q1(N) :- \text{Person}(N,L).$

- $Q \rightarrow Q1$:
 - CM: $N \rightarrow N, L \rightarrow L, C \rightarrow L, Z \rightarrow L, S \rightarrow L$
 - * $\text{CM}(\text{Person}(N,L)) = \text{Person}(N,L)$
 - * $\text{CM}(\text{Address}(L,C,Z,S)) = \text{Address}(L,L,L,L)$
- $Q1 \rightarrow Q$:
 - CM: $N \rightarrow N, L \rightarrow L$
 - * $\text{CM}(\text{Person}(N,L)) = \text{Person}(N,L)$

(d) Query Optimization with Constraints

- minimization of queries
 - remove body atoms (DL)
- find smallest query Q' such that $\text{body}(Q') \subseteq \text{body}(Q)$ and that $Q \equiv Q'$ given Σ

4. **TODO** more provenance

5. **TODO** GProM (just how to run it)

- source code
- on debussy: `/home/perm/semantic_opt_gprom`
- `./src/command_line/gprom -backend postgres -host 127.0.0.1 -user postgres -passwd test -port 5450 -db gpromtest -frontend dl`
- `-Osemantic_opt TRUE -Oflatten_dl TRUE`
- sqlite on: `./src/command_line/gprom -backend sqlite -db ./examples/test.db -frontend dl`
- time one query: `./src/command_line/gprom -backend postgres -host 127.0.0.1 -user postgres -db semanticopt -port 5433 -passwd test -frontend dl -timing -query 'Q(X) :- "r"(X,Y).'`

(a) compute lineage

$Q(X) :- R(X,Y), S(Y,Z).$ **ANS: Q. RP(1). FD R: A -> B. LINEAGE FOR R FOR**
 \hookrightarrow **RESULTS FROM RP.**

- time it query: `./src/command_line/gprom -backend postgres -host 127.0.0.1 -user postgres -db semanticopt -port 5433 -passwd test -frontend dl -Osemantic_opt TRUE -Oflatten_dl TRUE -timing -query 'Q(X) :- "r"(X,Y), "s"(Y,Z).` **ANS: Q. RP(1). FD "r": "a" -> "b". LINEAGE FOR "r" FOR RESULTS FROM RP.'**
- run query multiple times and time each execution: `./src/command_line/gprom -backend postgres -host 127.0.0.1 -user postgres -db semanticopt -port 5433 -passwd test -frontend dl -Osemantic_opt TRUE -Oflatten_dl TRUE -timing -time_queries TRUE -repeat_query_count 10 -query 'Q(X) :- "r"(X,Y), "s"(Y,Z).` **ANS: Q. RP(1). FD "r": "a" -> "b". FD "r": "b" -> "a". LINEAGE FOR "r" FOR RESULTS FROM RP.**

for postgres for now:

```
Q(X) :- "r"(X,Y), "s"(Y,Z). ANS: Q. RP(1). FD "r": "a" -> "b". LINEAGE
↳ FOR "r" FOR RESULTS FROM RP.
```

- just run a query and also time it (using SQLite int his example)
- `./src/command_line/gprom -backend sqlite -db ./examples/test.db -Osemantic_opt TRUE -Ofatten_dl TRUE -timing TRUE -loglevel 3 -frontend dl -query 'Q(X) :- R(X,Y), S(Y,Z). ANS: Q. RP(1). FD R: A -> B. LINEAGE FOR R FOR RESULTS FROM RP.'`
- **ANS:** the result relation for the query
- **FDs:** FD table: columns -> columns.
- **LINEAGE FOR R** - compute lineage of input table R
 - ... FOR RESULTS FROM RP. - then only compute lineage for results from RP

6. **TODO** how to run PDQ

7. **TODO** git

- <https://github.com/IITDBGroup/GProM>

8. practice CS520

(a) Write a Datalog program that returns the lastname and gpa of students that study cs

```
Q(lastname,gpa) :- student(_,_,lastname,major,gpa), major="cs"
Q(X,Y) :- student(S1, S2, X, "cs", Y)
```

(b) Surfing or hacking

```
Q(X,Y) :- student(S1, X, Y, S2, S3), interest(S1, "surfing").
Q(fname,lname) :- student(sid,fname,lname,_,_), interest(sid,"surfing").
```

- surfing or hacking:

```
Q(fname,lname) :- student(sid,fname,lname,_,_), interest(sid,A), (A =
↳ "surfing" OR A = "hacking").
```

- surfing and hacking:

```
Q(fname,lname) :- student(sid,fname,lname,_,_), interest(sid,"surfing"),
↳ interest(sid, "hacking").
```

- surfing or hacking:

```
Q(fname,lname) :- student(sid,fname,lname,_,_), interest(sid,A), A =
↳ "surfing".
```

```
Q(fname,lname) :- student(sid,fname,lname,_,_), interest(sid,A), A =
↳ "hacking".
```

$\forall sid, fname, lname, m, g, A : student(sid, fname, lname, m, g) \wedge interest(sid, A) \wedge A = "surfing" \rightarrow \exists Q$

$\forall fname, lname : Q(fname, lname) \rightarrow \exists sid, m, g, A : student(sid, fname, lname, m, g) \wedge interest(sid, A)$

(c) Students with same interest

```
Q(S1,L1,S2,L2):- Student(S1,f,L1,m,g), Interest(S1,a1),
↳ Student(S2,f,L1,m,g), Interest(S2,a2), a1 = a2, S1 < S2.
```

1.2.3 DONE terminal basics

1. DONE ssh

- for windows users use `putty` or WSL
- create terminal session on a different machine
 - connect as `USER`
 - to machine `MACHINE`

`ssh USER@MACHINE`

- copy files between machines `-scp`

`scp file otherfile`

2. DONE running programs

- program: run by inputting name
- options: pass after the program typically start with `-`
- `find dir options` - searching files in `dir`
- `man program` - open help for program
 - `SPACE` next page
 - `p` previous page
 - `q` quit
 - `/term` search for `term`
 - * `/<enter>` move to next match
- `cat file` - print file content
- `grep` search for content in files
- `echo msg` print `msg` to stdout

(a) combining programs

- input / output streams
 - `stdout`
 - `stderr`
 - = `stdin`
- `p1 | p2` - pass output of `p1` into `p2` (connect `p1` stdout to `p2` stdin)
- redirect `>` redirect stdout, `2>` redirect stderr
- read `stdin` from file with `< file`

3. DONE navigating directories

(a) directories & files

i. directories

- separated by `/`
- home directory `~`

- list content of current directory `ls`
 - `-a` list hidden files also
 - `-l` list file details
- `pwd` - prints the current directory
- `cd` - move to a different directory
 - `..` means one level up
 - `.` means the current folder
 - starting with `/` means absolute
 - without prefix `/` means relative to current directory
- ii. permissions
 - `r` - reading
 - `w` - writing
 - `x` - executable (files), can change into for directories
 - permission 9 values (3 for user (owner), 3 for group (owning), 3 for public (everybody else))
 - `chown user file` - change owner of file to user
 - `chgrp grp file` - change owner group of file to grp
 - `chmod permission file` - change permissions of file to permission
 - as 3 numbers (user, group, public) 4 means reading, 2 writing, 1 executing. Sum up these numbers
 - super power user: `root` can do everything
 - temporarily become root: `sudo`
- iii. deletion
 - `rm file` deletes file

1.2.4 TODO programming languages

1. TODO Python
2. TODO C
3. TODO Java

1.3 setup access to machines

1.3.1 machines

- `debussy.cs.iit.edu`

1.3.2 ssh

1. Mac
 - open terminal
 - run ssh with `user` and machine `machine`

`ssh user@machine`

for instance

```
ssh perm@debussy.cs.iit.edu
```

2. windows

- option 1: install Putty
- WSL -> like mac

1.3.3 running postgres

- **postgres cluster**: where the data is stored
- **postgres server program**: postgres or postmaster
 - -D tells postgres where the data will be stored
 - -c where to find the configuration file
 - -p which network port to run on
- **psql**: psql - run queries
 - -U USER - connect as user USER
 - -p PORT - port
 - -h HOST - host: 127.0.0.1
 - -D DATABASE - the database to connect to

1. psql

```
psql -h 127.0.0.1 -U postgres -p 5433 -D DBNAME
```

- \q - quit psql
- \? - help for all backslash commands
- \d OBJECT - print information about OBJECT (e.g., a table)
- to load data (by running a sql script)

```
psql# \i file
```

- TPC-H loading scripts /local/perm/tpchdata/scripts/ddl_1.sql - is 1GB
- creating database

```
CREATE DATABASE name;
```

2. postgres 10

- check with servers are running

```
ps aux | grep postgres
```

```
/usr/lib/postgresql/10/bin/postgres -D /var/lib/postgresql/10/main -c  
↪ config_file=/etc/postgresql/10/main/postgresql.conf
```

- connect to server

```
psql -h 127.0.0.1 -U postgres -p 5433 postgres
```

- connect with gprom

```
./src/command_line/gprom -backend postgres -host 127.0.0.1 -user postgres -db  
↪ semanticopt -port 5433 -passwd test
```

3. postgres 11

- connect to server

```
psql -h 127.0.0.1 -U postgres -p 5453 semanticopt
```

1.4 code and what to compare

1.5 experiments

1.5.1 the problem setting

- **input:** user query for provenance wrt. to query Q , database D , set of constraints Σ , input table R to a query result subset $R' \subseteq Q(D)$
- **step 1:** Generate query QP that computes provenance of R' in R for Q, D
- **step 2:** optimize the query to minimize it's size (to generate QP' equivalent to QP under a given set of constraints Σ)
- **step 3:** run the optimized query $QP'(D)$
- **output:** provenance which is subset of R

1.5.2 what parameters to vary?

- database size D
 - data distribution / real world or benchmark datasets
- structure and size of query Q
 - how selective is the query in terms of provenance
- number of constraints Σ
- which input table R
- what subset of results (R')

1.5.3 competitors

- **what methods to compare:**
 - **baseline:** do not optimize the query (free step 2, we pay at step 3)
 - **PDQ:** has an expensive step 2, but may be better sometimes in step 3 (complete method)
 - **our approach:** less expensive step 2, but may be worse in step 3 than PDQ

1.5.4 running step 2

```
for x in `seq 100`;
do
    ./src/command_line/gprom -backend postgres -host 127.0.0.1 -user postgres
    ↪ -db semanticopt -port 5453 -passwd test -frontend dl -Osemantic_opt TRUE
    ↪ -Oflatten_dl TRUE -loglevel 0 -Pexecutor sql -timing -queryFile
    ↪ ./umflint/tpcq18/customer.sql;
done \
| grep 'timer: TOTAL' \
| awk ' { print $5 }' \
> q18-opttime-customer.csv

rm q18-opttime-customer.csv; \
for x in `seq 100`;
do
    ./src/command_line/gprom -backend postgres -host 127.0.0.1 -user postgres
    ↪ -db semanticopt -port 5453 -passwd test -frontend dl -Osemantic_opt TRUE
    ↪ -Oflatten_dl TRUE -loglevel 0 -Pexecutor sql -timing -queryFile
    ↪ ./umflint/tpcq18/customer.sql \
| grep 'timer: TOTAL' \
| awk ' { print $5 }' \
>> q18-opttime-customer.csv
done
```

1.5.5 running step 3

1. generate provenance capture SQL queries

- generate file with optimized SQL query capturing provenance

```
./src/command_line/gprom -backend postgres -host 127.0.0.1 -user postgres -db
↪ semanticopt -port 5453 -passwd test -frontend dl -Osemantic_opt TRUE
↪ -Oflatten_dl TRUE -loglevel 0 -Pexecutor sql -queryFile
↪ ./umflint/tpcq18/customer.sql \
> ./umflint/tpcq18/p_customer.sql
```

- generate file with unoptimized SQL query capturing provenance

```
./src/command_line/gprom -backend postgres -host 127.0.0.1 -user postgres -db
↪ semanticopt -port 5453 -passwd test -frontend dl -Osemantic_opt FALSE
↪ -Oflatten_dl TRUE -loglevel 0 -Pexecutor sql -queryFile
↪ ./umflint/tpcq18/customer.sql \
> ./umflint/tpcq18/p_customer-unopt.sql
```

2. evaluate provenance capture queries

- time with psql (optimized)


```

for x in `seq 1 1000`; do \
    psql -h 127.0.0.1 -U postgres -d semanticopt -p 5453 -o /dev/null -c
    → '\timing on' -f ./umflint/tpcq18/p_customer.sql | grep 'Time:' | awk '
    → { print $2 }'; \
done > exp_results/tpcq18/p_customer.csv

```

- time with psql (unoptimized)

```

for x in `seq 1 1000`; do \
    psql -h 127.0.0.1 -U postgres -d semanticopt -p 5453 -o /dev/null -c
    → '\timing on' -f ./umflint/tpcq18/p_customer-unopt.sql | grep 'Time:' |
    → awk ' { print $2 }'; \
done > exp_results/tpcq18/p_customer-unopt.csv

```

3. tmux

- create a terminal session that continues after you disconnect from your ssh session
- create tmux

tmux

- detach from session CTRL-b d
- attach to existing tmux session (if our session is 5)

tmux list-sessions

tmux a -t 5

- create new window: CTRL-b c
- rename a window: CTRL-b ,
- jump to window numbered n: CTRL-b n, e.g., CTRL-b 0
- delete window: CTRL-b &

4. check system load (is somebody else utilizing the system for heavy work)

```

htop # show process / CPU / memory utilization
sudo iotop # show disk utilization (read / write)

```

5. generate TPC-H Datalog

(a) TPC-H Q3

- translate query to datalog

```

SELECT l_orderkey, -- select this column
       sum(l_extendedprice*(1-l_discount)) as revenue,
       o_orderdate,
       o_shippriority
FROM customer c, orders o, lineitem l
WHERE o_orderdate < '1995-03-15'
      AND l_shipdate > '1995-03-15'

```

```

    AND c.c_mktsegment = 'BUILDING'
    AND c.c_custkey = o.o_custkey
    AND l.l_orderkey = o.o_orderkey
GROUP BY l_orderkey, o_orderdate, o_shippriority
SELECT l_orderkey, -- select this column
       sum(l_extendedprice*(1-l_discount)) as revenue,
       o_orderdate,
       o_shippriority
FROM customer c JOIN orders o ON (c.c_custkey = o.o_custkey) JOIN
↪ lineitem l ON (l.l_orderkey = o.o_orderkey)
WHERE o_orderdate < '1995-03-15'
      AND l_shipdate > '1995-03-15'
      AND c.c_mktsegment = 'BUILDING'
GROUP BY l_orderkey, o_orderdate, o_shippriority

```

- equivalent datalog
- if aggregation function in head, then non-aggregated variables are group-by

```

Q(l_ok, sum(l_ep * (1-l_d)), o_od, o_sp) :-
    customer(c_ck,c_n,c_a,c_nk,c_p,c_ab,c_ms,c_ct),
    orders(o_ok,o_ck,o_os,o_t,o_od,o_op,o_c,o_sp,o_ct),
    lineitem(l_ok,l_pk,l_sk,l_ln,l_q,l_ep,l_d,x,y,z,a,b,c,d,e,f),
    o_od < '1995-03-15',
    l_sd > '1995-03-15',
    c_ms = 'BUILDING',
    c_ck = o_ck,
    l_ok = o_ok.

```

- with reusing variables instead of equality comparisons

```

Q(o_ok, sum(l_ep * (1-l_d)), o_od, o_sp) :-
    customer(c_ck,c_n,c_a,c_nk,c_p,c_ab,c_ms,c_ct),
    orders(o_ok,c_ck,o_os,o_t,o_od,o_op,o_c,o_sp,o_ct),
    lineitem(o_ok,l_pk,l_sk,l_ln,l_q,l_ep,l_d,x,y,z,a,b,c,d,e,f),
    o_od < '1995-03-15',
    l_sd > '1995-03-15',
    c_ms = 'BUILDING'.

```

- compute provenance for table customer

```

Q(o_ok, sum(l_ep * (1-l_d)), o_od, o_sp) :-
↪ customer(c_ck,c_n,c_a,c_nk,c_p,c_ab,c_ms,c_ct),
↪ orders(o_ok,c_ck,o_os,o_t,o_od,o_op,o_c,o_sp,o_ct),
↪ lineitem(o_ok,l_pk,l_sk,l_ln,l_q,l_ep,l_d,x,y,z,a,b,c,d,e,f), o_od
↪ < '1995-03-15', l_sd > '1995-03-15', c_ms = 'BUILDING'.

```

ANS : Q.

LINEAGE FOR customer.

- compute provenance for subset of results

```

Q(o_ok, sum(l_ep * (1-l_d)), o_od, o_sp) :-
  ⇨ customer(c_ck,c_n,c_a,c_nk,c_p,c_ab,c_ms,c_ct),
  ⇨ orders(o_ok,c_ck,o_os,o_t,o_od,o_op,o_c,o_sp,o_ct),
  ⇨ lineitem(o_ok,l_pk,l_sk,l_ln,l_q,l_ep,l_d,x,y,z,a,b,c,d,e,f), o_od
  ⇨ < '1995-03-15', l_sd > '1995-03-15', c_ms = 'BUILDING'.

```

ANS : Q.

```
QP(a,b,c,d) :- Q(a,b,c,d), a = 1231455.
```

```
LINEAGE FOR customer FOR RESULTS FROM QP.
```

1.5.6 TODO running things in PDQ

1.5.7 experiment dry run

1.6 TPC-H datalog queries

1.6.1 Q3

1. capture provenance for customer

```

q(l_ok, sum(l_ep*(1-l_d)), o_od, o_sp) :-
customer(c_ck,c_n,c_a,c_nk,c_p,c_ab,'BUILDING',c_ct),
orders(l_ok,c_ck,o_os,o_t,o_od,o_op,o_c,o_sp,o_ct),
lineitem(l_ok,l_pk,l_sk,l_ln,l_q,l_ep,l_d,x,y,z,l_sd,b,c,d,e,f),
o_od < '1995-03-15', l_sd > '1995-03-15'.

```

ANS: q.

```
LINEAGE FOR customers FOR RESULTS FROM q.
```

2. capture provenance for lineitems

```

q(l_ok, sum(l_ep*(1-l_d)), o_od, o_sp) :-
customer(c_ck,c_n,c_a,c_nk,c_p,c_ab,'BUILDING',c_ct),
orders(l_ok,c_ck,o_os,o_t,o_od,o_op,o_c,o_sp,o_ct),
lineitem(l_ok,l_pk,l_sk,l_ln,l_q,l_ep,l_d,x,y,z,l_sd,b,c,d,e,f),
o_od < '1995-03-15', l_sd > '1995-03-15'.

```

ANS: q.

```
LINEAGE FOR lineitems FOR RESULTS FROM q.
```

3. capture provenance for orders

```

q(l_ok, sum(l_ep*(1-l_d)), o_od, o_sp) :-
customer(c_ck,c_n,c_a,c_nk,c_p,c_ab,'BUILDING',c_ct),
orders(l_ok,c_ck,o_os,o_t,o_od,o_op,o_c,o_sp,o_ct),
lineitem(l_ok,l_pk,l_sk,l_ln,l_q,l_ep,l_d,x,y,z,l_sd,b,c,d,e,f),
o_od < '1995-03-15', l_sd > '1995-03-15'.

```

ANS: q.

LINEAGE FOR orders FOR RESULTS FROM q.

1.6.2 Q18

2 Meeting Murali

Added: [2021-11-12 Fri 13:46]

- schema: R(A,B), S(C,D)

Q(X) :- R(X,Y), S(Y,Z).

$QP \subseteq Q$

PROV_R(X,Y) :- R(X,Y), S(Y,Z). QP(X).

functional dependency: A -> B

PROV_R(X,Y) :- R(X,Y). QP(X).

A	B	C
1	1	a
2	2	c
2	2	d

$\forall x, y, z, x', z' : address(x, y, z) \wedge address(x', y, z') \rightarrow z = z'$
zip -> city

street	zip	city
10	60614	NY
23	60614	NY